

AUTOMATING HPC SIMULATIONS ON RESCALE CLOUD PLATFORM USING REST API

Karthik Reddy Lyathakula

PhD Candidate, North Carolina State University

GitHub: <https://github.com/karthikncsu/Rescale-Tutorial>

CONTENTS


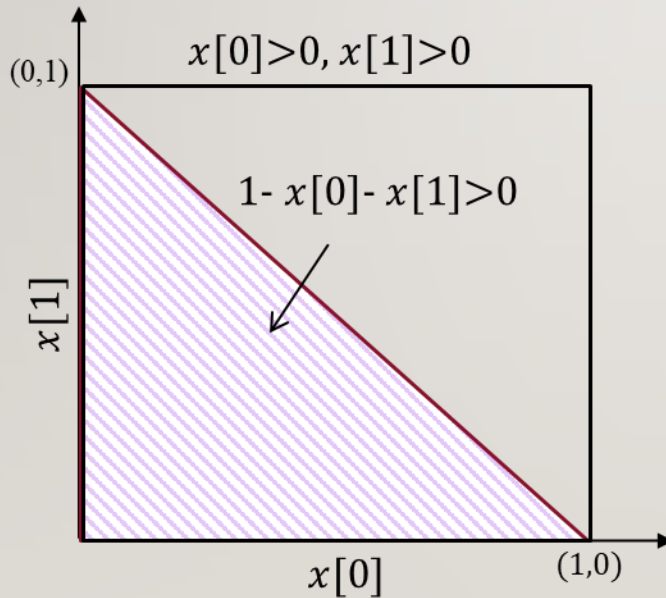
- Problem: Sampling from High Dimensional Space
- Sampling Methods
- Parallelization of Sampling
 - Sequential Monte Carlo (SMC) Method
- Benchmarking Speed-up of SMC code
 - Simulation on Rescale Cloud Platform
 - Automating multi jobs using REST API

SAMPLING FROM HIGH DIMENSIONAL SPACE

n -dimensional space with K non-linear constraints

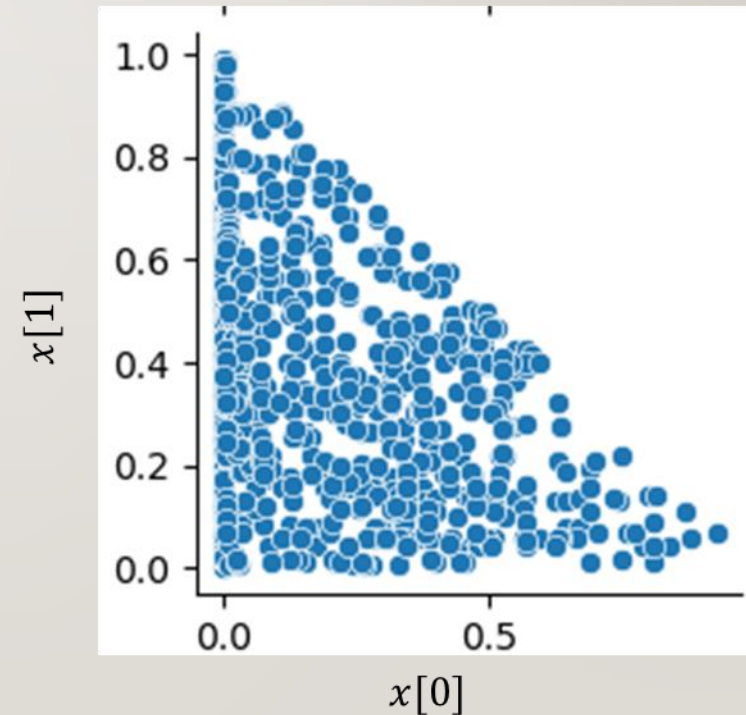
$$C_k(\mathbf{x}) > 0, k = 1, \dots, K$$

$$\mathbf{x} = (x_1, \dots, x_n)$$



- Used in many applications
 - Uncertainty quantification
 - Inverse problems

Generated Samples



SAMPLING METHODS

- Metropolis-Hastings Random Walk
- Adaptive Metropolis
- Gibbs Sampler
- Sequential Monte Carlo (SMC)





Can be parallelized

MCMC methods

Generates samples sequentially

● Accepted

● Rejected



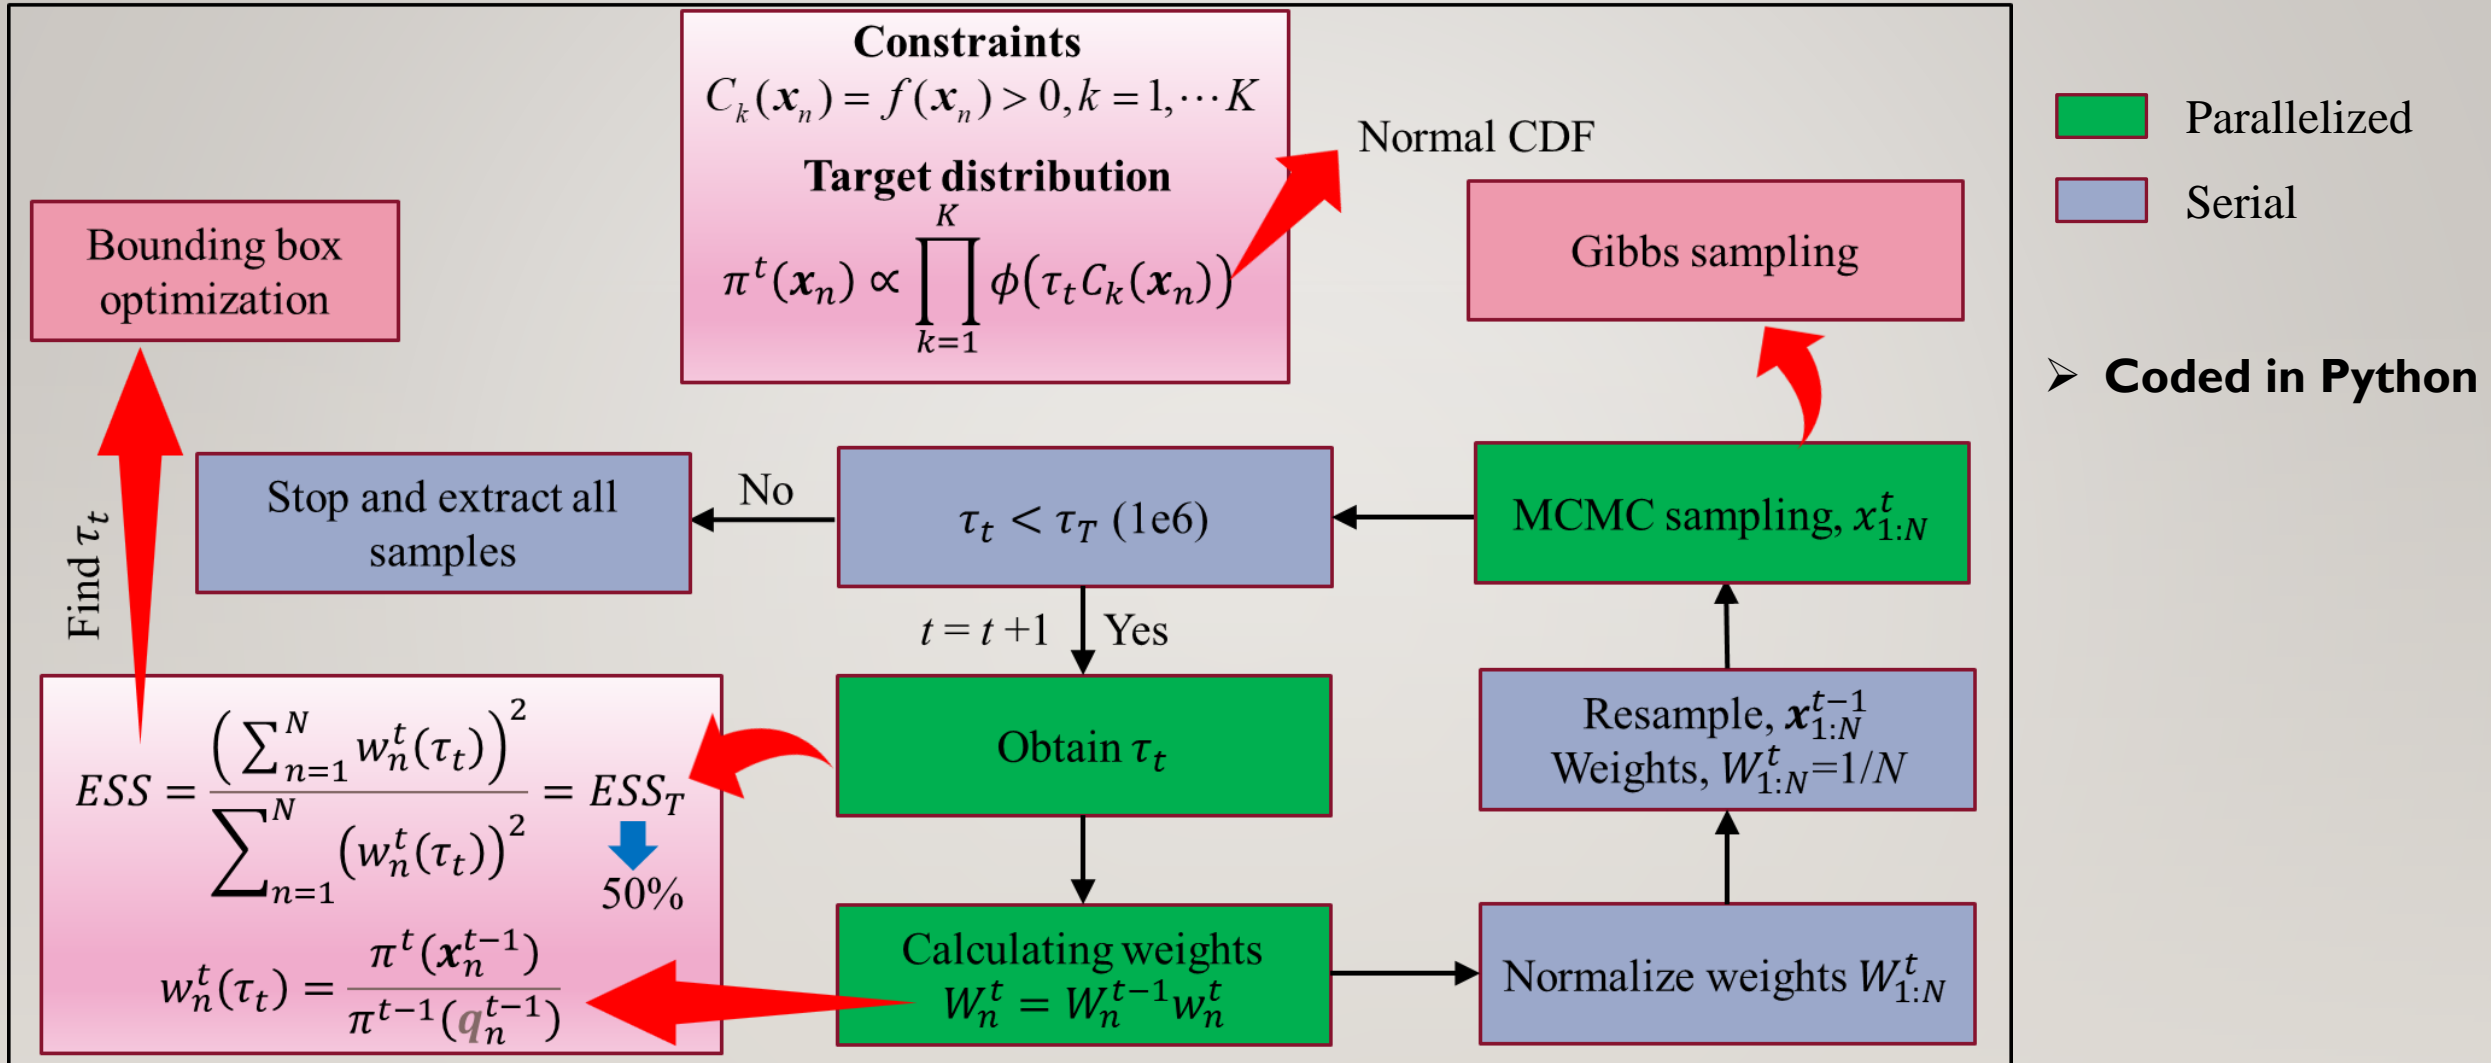
MCMC Methods: Fails to generate samples

SMC Methods: Can generate samples

High Dimensional Space

$$\begin{aligned} 1.0 - x[0] - x[1] - x[2] - x[3] &\geq 0.0 \\ x[3] / (x[0] + x[1]) - 0.1 &\geq 0.0 \\ 0.15 - x[3] / (x[0] + x[1]) &\geq 0.0 \\ x[2] - 0.5 &\geq 0.0 \\ x[0] + x[1] / (x[0] + x[1] + x[3]) - 0.8 &\geq 0.0 \end{aligned}$$

SEQUENTIAL MONTE CARLO (SMC)



GitHub: <https://github.com/karthikncsu/Sampling-from-high-dimensional-space>

BENCHMARKING SMC CODE

- Benchmarking parallelization of SMC code
 - Calculating speedup (S_p) and efficiency (E_p)

$$S_p = \frac{T_s}{T_p} \quad E_p = \frac{T_s}{nT_p}$$

T_s : Sequential simulation time

T_p : Parallelized simulation time

Simulation on Rescale cloud platform

- Multiple SMC simulations
 - Using 1, 2, 4, 8, 18 processors
- Automated simulations
 - Submit multi jobs using REST API
 - Download SMC log files once completed
 - Plot speedup and efficiency

AUTOMATING MULTI JOBS USING REST API

- Python is used for REST API
- Implemented using python class: BatchJobSubmit
- Steps for automating simulation
 - Step 1 : Upload input file
 - Step 2 : Setup and Submit job
 - Step 3 : Monitor status of Jobs
 - Step 4 : Download SMC log files
 - Step 5 : Postprocessing

AUTOMATING MULTI JOBS USING REST API

```
def file_upload(self, input_file, oncloud=False):  
    """  
    Function for uploading input files  
    """  
    if oncloud:  
        self.file_ids.append(file_id)  
        return True, file_id  
    else:  
        #Uploading a file to teh rescale cloud  
        file_upload = requests.post(  
            'https://platform.rescale.com/api/v2/files/contents/',  
            data=None,  
            files={'file': open(input_file, 'rb')},  
            headers={'Authorization': self.myapi_token}  
        )  
  
        status=self.print_status(file_upload, "File upload status")  
        if status:  
            upload_details=json.loads(file_upload.text)  
            file_id=upload_details["id"]  
            print("File ID:", file_id)  
            self.file_ids.append(file_id)  
            return status  
        else:  
            return status
```

Step I: Upload input file

- Place entire code into single folder
- Create a zip file
- File id is created for successfully uploaded input files

AUTOMATING MULTI JOBS USING REST API

```
def setup_submit_job(self,data):
```

```
    """
    Function for setting up and submitting the job
    """
```

```
    job_setup=requests.post(
        'https://platform.rescale.com/api/v2/jobs/',
        json=data,
        headers={'Content-Type': 'application/json',
                 'Authorization': self.myapi_token}
    )
```

```
    status=self.print_status(job_setup,"Job setup status")
```

```
    if status:
```

```
        job_details=json.loads(job_setup.text)
```

```
        job_id=job_details["id"]
```

```
        print("Job ID:",job_id)
```

```
        job_submit=requests.post(
            'https://platform.rescale.com/api/v2/jobs/%s/submit/'%job_id,
            headers={'Content-Type': 'application/json',
                     'Authorization': self.myapi_token}
        )
```

```
        status=self.print_status(job_submit,"Job submitting status")
```

```
    if status:
```

```
        self.job_ids.append(job_id)
```

```
    else:
```

```
        self.job_ids.append(None)
```

```
    return status
```

Step 2 : Setup and Submit job

➤ Select

- Number of processors
- Core type
- Input file id
- Commands to be executed
- Other details.....

```
def jobdata(file_id,numproc):
```

```
    command="cd Sampling-from-high-dimensional-space\n"
```

```
    command=command+"cd code\n"
```

```
    command=command+'conda install -y mpi4py \nmpirun -np '+str(numproc)+' python sampler.py alloy.txt alloy.out 20000 SMC'
```

```
    data={
        'name': 'Intern assignment sampling '+str(numproc),
        'projectId': 'gxzXS',
        'jobanalyses': [
            {
                'useMpi': True,
                'command': command,
                'analysis': {
                    'code': 'anaconda',
                    'version': '5.3.1'
                },
                'hardware': {
                    'coresPerSlot': numproc,
                    'slots': 1,
                    'coreType': 'hpc-3'
                },
                'inputFiles': [
                    {
                        'id': file_id
                    }
                ]
            }
        ]
    }
```

AUTOMATING MULTI JOBS USING REST API

Step 2 : Setup and Submit job

- Multiple job submissions using loop

```
input_file='Sampling-from-high-dimensional-space.zip'

jobsobj=BatchJobSubmit()
status_input=jobsobj.file_upload(input_file)
▼ if status_input:
▼     for numproc in [1,2,4,8,18,36]:
        data=jobdata(jobsobj.file_ids[0],numproc=numproc)
        status_submit=jobsobj.setup_submit_job(data)

jobsobj.status_job()
```

AUTOMATING MULTI JOBS USING REST API

```
def status_ind_job(self):
    """
    Function for status of all submitted jobs
    """

    job_statuses=[]
    for job_id in self.job_ids:
        if job_id is not None:

            job_status=requests.get(
                'https://platform.rescale.com/api/v2/jobs/%s/statuses/'%job_id,
                headers={'Content-Type': 'application/json',
                        'Authorization': self.myapi_token})

            output_json=json.loads(job_status.text)
            status=output_json["results"][0]["status"]
            job_statuses.append(status)

            statusReason=output_json["results"][0]["statusReason"]
            print("Job Id:",job_id," Status:",job_statuses[-1]," StatusReason:",statusReason)

        else:
            job_statuses.append("Completed")

    return job_statuses

def status_job(self):
    """
    Output for status of all the jobs
    """
    while True:
        time.sleep(15)

        print("-----",datetime.now(),"-----")
        if len(self.job_ids)==0:
            break
        job_statuses=self.status_ind_job()
        if job_statuses.count('Completed')==len(job_statuses):
            break
```

Step 3 : Monitor status of Jobs

```
✓ TERMINAL

Job Id: jewkhhb , Status: Completed , StatusReason: Completed successfully
Job Id: ZTVkhhb , Status: Completed , StatusReason: A run failed
----- 2021-03-29 13:21:09.618741 -----
Job Id: deWkhhb , Status: Executing , StatusReason: None
Job Id: JrczW , Status: Executing , StatusReason: None
Job Id: XTVkhhb , Status: Executing , StatusReason: None
Job Id: VBczW , Status: Completed , StatusReason: Completed successfully
Job Id: jewkhhb , Status: Completed , StatusReason: Completed successfully
Job Id: ZTVkhhb , Status: Completed , StatusReason: A run failed
----- 2021-03-29 13:21:25.426903 -----
Job Id: deWkhhb , Status: Executing , StatusReason: None
Job Id: JrczW , Status: Executing , StatusReason: None
Job Id: XTVkhhb , Status: Executing , StatusReason: None
Job Id: VBczW , Status: Completed , StatusReason: Completed successfully
Job Id: jewkhhb , Status: Completed , StatusReason: Completed successfully
Job Id: ZTVkhhb , Status: Completed , StatusReason: A run failed
----- 2021-03-29 13:21:41.318903 -----
Job Id: deWkhhb , Status: Executing , StatusReason: None
Job Id: JrczW , Status: Executing , StatusReason: None
Job Id: XTVkhhb , Status: Executing , StatusReason: None
Job Id: VBczW , Status: Completed , StatusReason: Completed successfully
Job Id: jewkhhb , Status: Completed , StatusReason: Completed successfully
Job Id: ZTVkhhb , Status: Completed , StatusReason: A run failed
```

- Job with 36 processors failed due to technical issue with Rescale
 - mpi4py module cannot be installed using 36 processors

AUTOMATING MULTI JOBS USING REST API

```
def download_file(self,filename,job_id):
    """
    Download files from job numbers
    """

    #Finding all the url to download the files
    url='https://platform.rescale.com/api/v2/jobs/%s/files/%s/job_id'
    headers={'Authorization': self.myapi_token}

    file_download=requests.get(url, headers=headers)
    status=self.print_status(file_download,"File download status")
    file_found=False

    if status:
        job_details=json.loads(file_download.text)

        while True:
            #Get all the details of the files in the current page
            for arg in job_details["results"]:
                if arg["name"]==filename:
                    file_found=True
                    downloadUrl=arg["downloadUrl"]

                    filename_save="./downloads/out_"+job_id+"_"+filename
                    fout=open(filename_save,"wb")
                    with closing(requests.get(downloadUrl, headers=headers, stream=True)) as r:
                        with open(filename, 'r') as f:
                            for chunk in r.iter_content(chunk_size=100):
                                fout.write(chunk)
                    fout.close()
                    fout=open(filename_save,"r")
                    data=fout.readlines()
                    fout.close()
                    return True,data

            #Checking for details of the files in the next page
            if job_details["next"]==None:
                break
            url=job_details["next"]

            file_download=requests.get(url, headers=headers)
            job_details=json.loads(file_download.text)

        return False,None
```

Step 4 : Download SMC log files

```
all_proc_data=[]
for job_id in jobsobj.job_ids:
    status,data=jobsobj.download_file("comp_time.log",job_id)
    if status:
        all_proc_data.append(data[0].split())
```

- Computational time is saved in comp_time.log by SMC on Rescale cloud

AUTOMATING MULTI JOBS USING REST API

```
print(all_proc_data)
if len(all_proc_data)>0:
    all_proc_data=np.asarray(all_proc_data).astype(np.float64)

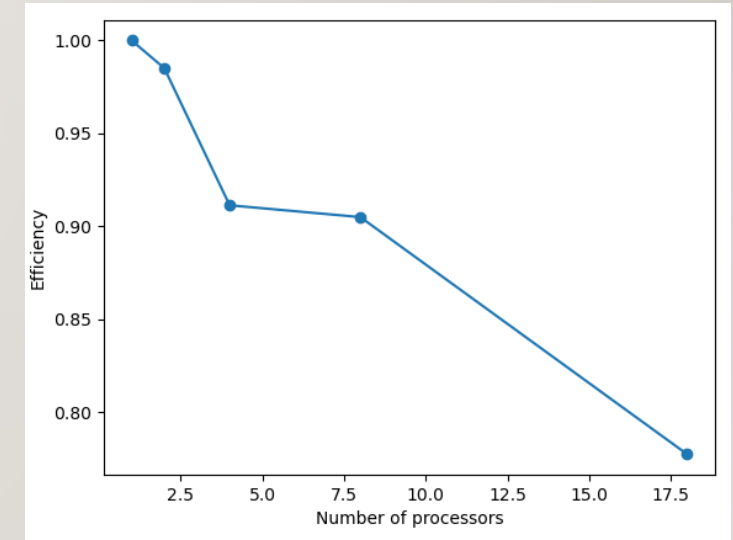
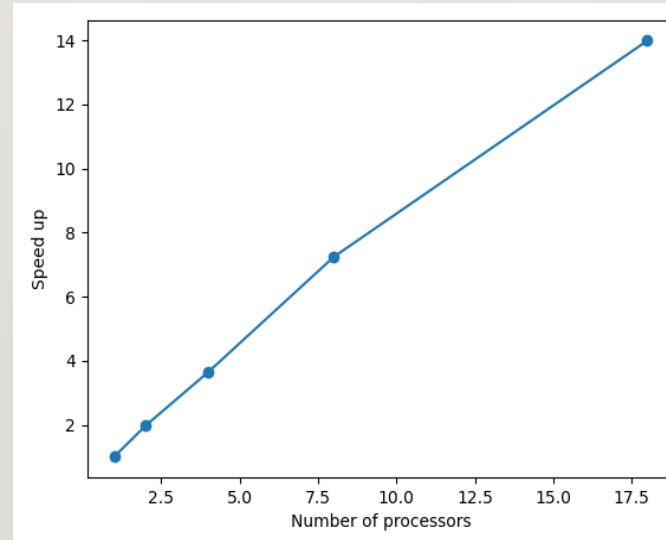
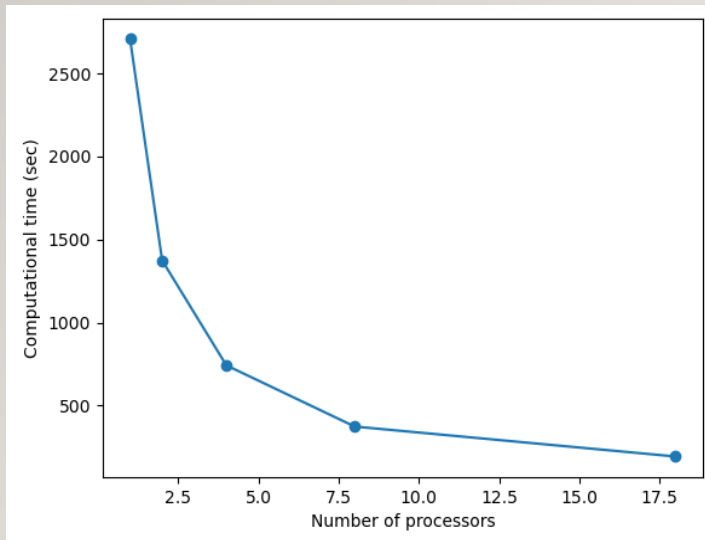
    fig=plt.figure()
    plt.plot(all_proc_data[:,0],all_proc_data[0,1]/all_proc_data[:,1],"o-")
    plt.xlabel("Number of processors")
    plt.ylabel("Speed up")
    fig.savefig("speedup.png")
    plt.close()

    fig=plt.figure()
    plt.plot(all_proc_data[:,0],all_proc_data[0,1]/(all_proc_data[:,1]*all_proc_data[:,0]),"o-")
    plt.xlabel("Number of processors")
    plt.ylabel("Efficiency")
    fig.savefig("efficiency.png")
    plt.close()

    fig=plt.figure()
    plt.plot(all_proc_data[:,0],all_proc_data[:,1],"o-")
    plt.xlabel("Number of processors")
    plt.ylabel("Computational time (sec)")
    fig.savefig("comptime.png")
    plt.close()
```

Step 5 : Postprocessing

- Computational time decreased with number of processors
- Linear speed up is achieved



THANK YOU

